

Context Window

When you type 2 + 2 on the calculator on your phone, you always get 4. Same input, same answer, every time. AI doesn't work that way. Consider the following: identical questions, asked to AI by two different people.



Luke

Asks AI

What car should I buy after I graduate from college?

ChatGPT answers

Great question, Luke. I definitely recommend a **Jeep Cherokee**.



Nate

Asks AI

What car should I buy after I graduate from college?

ChatGPT answers

I'm happy you asked this question, Nate. At this point in your life, I recommend the **Ford Raptor**.

Wait a second. In the last lab, you changed the answer by changing the prompt. But this time, Luke and Nate typed the exact same prompt into the same app, and still got different answers!

Some people think this is a reason not to trust AI. But the opposite is true. It's actually demonstrating the power of AI. It gave a different, more tailored answer to each.

That's the big idea of this lesson: **with the right context, AI gives better answers**. The rest of this page is about where that context comes from, and how to take control of it.

WHY DID THIS HAPPEN?

ChatGPT sees more than just the last message Luke and Nate typed. For example, earlier in his chat, Nate said he loves pickup trucks. When ChatGPT answers his question, it recognizes this and suggests the Ford Raptor.

What do you call everything the model can see when it answers your question? The **context window**. Think of it almost like the working memory the model uses to answer. And it builds that working memory from five places.



WHAT FOLLOWS YOU WHERE

Those five pieces don't all behave the same way. Each has a different reach: how far it follows you as you open new chats. Two live and die with a single chat. Two follow you everywhere. (The fifth, Projects, gets its own section below.)

STAYS IN THIS CHAT



Your current prompt

Everything you sent this turn: what you typed, plus anything attached.



Everything earlier in this chat

Everything you and the AI said before now.

FOLLOWS YOU EVERYWHERE



Personalization

What you told it once: your name, what you're into, how you like answers.



Saved Memory

What the app picked up from your chats and saved on its own.

PERSONALIZATION AND SAVED MEMORY

Personalization is the Luke and Nate effect, on purpose. Tell the app once what to call you, what you're into, and how you like your answers, and every chat starts already knowing it. Set it up in two minutes, and a generic tool becomes yours.

Saved Memory is the app noticing things in your chats worth keeping. It keeps a record of them on its own. That can sound like the AI learning about you, but remember: the model learned once, during training, and nothing you type changes it. It's the app keeping notes, and putting those notes into the window where the model sees them.

WHAT'S A PROJECT?

That leaves the dashed card in the illustration: Projects. A project is a folder for one piece of ongoing work: your summer job hunt, your team's season, the game you're building. You put that work's instructions and files inside: your resume, the schedule, your code, your notes.

Here's the idea: the big, ongoing things you do with AI have their own context. Your job hunt has a resume and a list of places you've applied. Your team has a schedule and a playbook. Without a project, you rebuild that context in every chat: re-upload the files, re-explain the work. A project holds the context for you. Start a chat inside it, and you step into that context instead of rebuilding it.

THE FORGETTING PROBLEM

Ever been deep in a long chat, and AI seems to forget something you told it earlier? There's a reason. The window isn't infinite, and when a chat gets long enough, the oldest parts fall out. Some apps summarize them, others just drop them. That's one reason starting a fresh chat for a new task is often the right move.



With the right context, AI gives better answers.

Control the context, control the quality.