

Document Trap

Your season-end basketball tournament starts next week, so you want to brush up on the rules. You upload your league’s 200–page rulebook and ask: “How many fouls until I’m out of the game?”

ChatGPT answers: “Five fouls and you foul out.”

Wait a second. In last year’s tournament, you remember a player picking up five fouls and staying in the game. So you dig through the rulebook yourself. Regular season: five fouls, exactly what ChatGPT said. But near the end there’s a special section for tournaments, and in those games players get six.

The AI pulled the standard limit and missed the exception. The answer wasn’t made up. It was incomplete. **The trap is thinking ‘uploaded’ means ‘fully read.’**

IT’S ABOUT THE CONTEXT WINDOW

As you learned earlier in the course, the model reads its full context window every time you chat. Your two-hundred-page rulebook comes out around 100,000 tokens. That’s more than some models will load at once, and even when it fits, the system still needs room for your conversation and its answer.

So when a long document doesn’t fit, and many don’t, the system runs a process on it instead:

How Chunks Get Picked

1 Your document gets split into chunks

Before you ever ask a question, the rulebook gets cut into smaller pieces, each a paragraph or two long. A few hundred chunks total. The chunks sit in a database, waiting.

2 Each chunk gets a meaning vector

Same idea you saw in the Embeddings lesson. Every chunk gets translated into a meaning profile, a list of numbers that captures what the chunk is about. The chunk about technical fouls ends up close in meaning to other rulebook chunks about penalties, referee calls, and player conduct.

3 Your question becomes a vector too. The closest chunks get loaded.

Now you ask: “Can a player call a timeout during a free throw?” Your question gets its own meaning vector. The system finds the chunks whose vectors are closest. Those chunks (not the whole rulebook) get loaded. The model answers based on what got retrieved.

WHAT GOT LOADED

The model never saw the whole rulebook.

It answered from only the chunks that matched your question. If the chunk that mattered didn’t get pulled, the answer still sounds confident. It is just missing what it never loaded.

There’s a name for what just happened: **retrieval**. When the system can’t load everything, it retrieves the pieces that match your question. Done well, this finds you a specific answer in a 200–page rulebook in seconds. Done poorly, the wrong pieces get pulled, and the model answers from incomplete evidence.

You can’t always tell what got retrieved. But you can steer it. Five moves help.

The five moves

1 Point to the section by name

Include keywords from the document in your question. The system uses them to find matching chunks.

VAGUE

What are the rules on flagrant fouls?

POINTED

What does the technical fouls chapter say about flagrant fouls?

2 Ask one question at a time

Multi-part questions force the system to retrieve different chunks for different parts. Break it up. One focused question per turn gives retrieval the best shot.

COMBO QUESTION

Summarize the policy, list the exceptions, and tell me what to do.

ONE THING AT A TIME

What does the policy say about late assignments? Quote the line.

3 Paste the section directly into your prompt

If you already know the paragraphs that matter, paste them into your message. That puts them in context with no retrieval lottery.

RETRIEVAL LOTTERY

Does my technical fouls clause apply if both players are involved?

PASTED IN DIRECTLY

Here’s the clause: ‘A technical foul is assessed when...’ Does this apply if both players are involved?

4 Upload only what’s relevant

Instead of the whole rulebook, upload just the chapter you need. Fewer chunks competing for retrieval slots means the right ones are more likely to make the cut.

WHOLE RULEBOOK

[Uploads the entire 200–page rulebook] What does it say about flagrant fouls?

JUST THE CHAPTER

[Uploads only the 3–page fouls chapter] What does it say about flagrant fouls?

5 Ask the AI to quote

Add ‘quote the exact section you’re basing this on’ to your question. If the quote is missing or doesn’t match the document, retrieval probably failed.

NO RECEIPTS

What does the rulebook say about flagrant fouls?

ASK FOR THE QUOTE

What does the rulebook say about flagrant fouls? Quote the exact line you’re basing it on.

All five share one idea: make the right chunks easy to find. The more specific your question, the more reliable the answer.



Uploaded does not mean understood. Confident does not mean complete. Asking for evidence is the skill.

It answers from what it retrieved.

Not from having read the whole thing.