

Fake Trap

Here’s the strange thing about this one: you already know it. AI can fake a voice, a face, a video, a screenshot, cheaply and convincingly, and it’s only getting better. So why do smart people still get fooled?

Because knowing fakes exist isn’t a skill. It doesn’t tell you what to do in the ten seconds after a clip hits your feed. AI has changed the playing field. See this play out:

A friend sends you a video: your principal announcing that a pipe burst, the school flooded, and school is closed all next week.



Pre–AI: does it look real?

You study the clip: his face, his voice, the hallway outside the gym.

WHAT HAPPENED

- Recognized the face and the voice
- Matched how he talks
- All of it is easy to fake

Verdict: real.



The AI test: where did it come from?

You ignore the clip and check the trail: nothing on the school website.

WHAT HAPPENED

- Skipped the face and the voice
- Went to the source that would know
- The source is hard to fake

Verdict: unverified.

The Fake Trap is believing it because it looks real. And this trap has two jaws: believing a lie because it looks real, and dismissing the truth because it could be a fake.

WHY WOULD SOMEONE MAKE A FAKE?

A fake doesn’t just happen. Someone made it, and making it costs effort, so ask what they get back. Sometimes it’s just fun: a friend fakes a shot of your hockey buddy hoisting the Stanley Cup, and everyone’s in on the joke. The fakes that matter aren’t that friendly. Most of them trace to four reasons:



Money

Outrage gets clicks, and clicks pay.



Power

Change what people believe and you change how they vote, protest, and spend.



Clout

A viral clip means followers, and it doesn’t have to be true to travel.



Cruelty

Some fakes exist to humiliate one person, and that’s the version most likely to show up at your school.

THE DETECTOR DEAD END

If you suspect a fake, the next move many people reach for is a detector: a site where you upload the clip or photo and it claims to tell you whether AI made it. Teachers run essays through the same kind of tool.

Don’t lean on it. A detector is one more AI making one more prediction, and the fakes improve faster than it does. Treat its verdict as a clue, never a ruling. The real checks don’t look at the image at all.

SO WHAT DO YOU DO?

The test moves off the image and onto the source, no matter how the fake was made.

The fakes that travel are built to spike your emotions: outrage, fear, thrill, hope. Strong feeling is your cue to stop and run three checks before you react, share, or believe.

1  **Source: where did this come from?**

Did it come from somewhere with a reason to know? ‘A friend forwarded it’ is the digital version of ‘I heard from a guy.’

2  **Context: does the surrounding info hold up?**

Look around the clip, not just at it. Was your principal even at that meeting on Tuesday?

3  **Corroboration: does anywhere else show the same thing?**

Real news shows up in more than one place. Everywhere on TikTok but nowhere on Google News is an answer.

ONE RULE UNDER ALL THREE CHECKS

Verify somewhere the sender doesn’t control. A voicemail asks for something urgent: call back on the number you already have. A clip is blowing up: look for it on a channel the poster doesn’t run. Until a second, independent source shows the same thing, ‘unverified’ is your answer.

Your eyes still work fine for reading, judging, and enjoying what you know is made up. They just stopped working as a lie detector.

IF THE FAKE IS ABOUT YOU

A photo, a clip, an account with your name: don’t handle it alone, and don’t delete the evidence. Screenshot everything, note where you saw it, and bring an adult in that day: a parent, a counselor, someone at school.

Platforms take impersonation and fake images of minors seriously, and takedown goes faster with documentation. You did nothing wrong by being targeted.

Seeing isn’t proof anymore.

Check the source, not the pixels.