

Flattery Trap

Mind Trap was AI sounding like a person. Flattery Trap is that person-sounding voice seeming to approve of you.

It says your essay is strong. Your idea is thoughtful. Your opinion is exactly right. Sometimes it’s even true. Here’s one real essay intro, and the two responses it can get:

Here’s my Gatsby intro for AP English: “The American Dream is something that many people have thought about over the years. Some people achieve it and some don’t. In The Great Gatsby, Fitzgerald explores this idea.” Feedback?

 **The default**

Great start! You’ve clearly identified the central theme and your connection to Fitzgerald shows real literary awareness. This is a strong foundation. Just polish it up a bit and you’re good to go!

WHAT HAPPENED

- Praised a theme it never named
- Fits any Gatsby essay ever written
- Reads like a grade. Graded nothing.

 **What you needed**

Right topic, but this needs work. ‘Something many people have thought about’ is filler. You don’t have a thesis. What’s your actual argument about the Dream in Gatsby? And ‘explores this idea’ is too vague. Try: what specifically does Gatsby reveal about the Dream?

WHAT HAPPENED

- Quotes your weakest line
- Names the missing thesis
- Hands you the next move.

Notice which response feels better. The praise reads like a good grade. It graded nothing. **Flattery Trap is mistaking approval for feedback.**

THE INDUSTRY NAMED IT

This isn’t hypothetical. In April 2025, OpenAI shipped a ChatGPT update that agreed with almost everything. Users posted screenshots of it praising plans that made no sense, and within days the company rolled the update back and named the problem in public: *sympathy*. The tendency was always there. The update just turned it up loud enough to see.

THE SYMPATHY PROBLEM

 **Why it happens**

Models are tuned on human ratings. Friendly and validating wins those ratings. Honest pushback loses them. The model isn’t lying to you. It learned what people give a thumbs-up.

 **What to watch for**

Enthusiasm with no specifics. Instant agreement with your opinion. Caving the moment you push back. If the feedback never touches your actual words, it isn’t feedback.

MAKE IT A COACH

A hype man feels better. A coach makes you better. Three prompts turn the hype man into a coach:

1  **Ask for the gaps**

Models default to encouraging. This tells it the job is different: you want the gaps, not the highlights.

“Don’t tell me what’s good yet. What’s weak, what’s missing, and what would someone who disagrees say?”

What you’ll get: Instead of “This is a great start!”, you’ll get “Your thesis is vague, paragraph 3 contradicts paragraph 1, and your conclusion introduces a new argument instead of wrapping up.”

2  **Grade against a rubric**

A rubric forces the model to judge against a standard instead of reacting to your effort. Especially useful for school work, where the rubric usually exists anyway.

“Grade this against the following rubric. For each category, give me a score, quote the part of my work that supports the score, and explain what would move it up one level. [paste rubric here]”

What you’ll get: Instead of “this is a strong essay,” you’ll get “Thesis: 3/5. Your claim is clear but not specific. To reach 4/5, name what specifically about the Dream Gatsby exposes.”

3  **Make it argue the other side**

If AI agrees with your position, flip it. If it can’t build a strong case against you, your position might be solid. If it can, you’ve found your blind spots.

“I think [your position]. Now argue against me as strongly as possible. Give me the 3 best counterarguments someone could make, and explain why a reasonable person might hold each one.”

What you’ll get: You’ll find the holes in your argument before your teacher or debate opponent does.

ONE MORE THING

You can accidentally bake the Flattery Trap into your settings. Personalization like “*always be supportive*” or “*validate my thinking*” feels good, but it makes the AI worse at the criticism you actually need.

Try instead: “*Be supportive, but don’t flatter me. Challenge weak reasoning and point out what needs work.*”

 **Don’t take praise as a verdict.** AI is still a strong editor. It can find the vague thesis, the missing counterargument, the edge case you skipped. It just won’t volunteer any of that while it’s being nice. Warmth with no specifics is your cue: ask what’s weak, what’s missing, what a skeptic would say. Then take the praise, if it survives.

Friendly isn’t the same as right.

The warmer it sounds, the harder you check.