

Flattery Trap

Mind Trap was AI sounding like a person. Flattery Trap is that person-sounding voice seeming to approve of you.

It says your essay is strong. Your idea is thoughtful. Your opinion is exactly right. **The Flattery Trap is mistaking AI’s approval for feedback.**

Sometimes it’s even true. You ask AI for feedback on your essay intro to The Great Gatsby.

The American Dream is something that many people have thought about over the years. Some people achieve it and some don’t. In The Great Gatsby, Fitzgerald explores this idea.

 **The Flattery Trap**

Great start! You’ve clearly identified the central theme and your connection to Fitzgerald shows real literary awareness. This is a strong foundation. Just polish it up a bit and you’re good to go!

WHAT HAPPENED

- Praised a theme it never named
- Fits any Gatsby essay ever written
- Reads like a grade. Graded nothing.

 **What you needed**

Right topic, but this needs work. The words ‘Something many people have thought about’ are filler. You don’t have a thesis. What’s your actual argument about the American Dream in The Great Gatsby? And ‘explores this idea’ is too vague. Try: what specifically does Gatsby reveal about the Dream?

WHAT HAPPENED

- Quotes your weakest line
- Names the missing thesis
- Hands you the next move.

Notice which response feels better. The praise reads like a good grade. It graded nothing.

THE INDUSTRY NAMED IT

This isn’t hypothetical. In April 2025, OpenAI shipped a ChatGPT update that agreed with almost everything. The screenshot that made it famous: a Reddit user pitched a gag product, poop on a stick, to see how far the flattery would go. ChatGPT called the idea genius, “performance art disguised as a gag gift,” and encouraged them to invest \$30,000.

Within days the company rolled the update back and named the problem in public: *sycophancy*. The tendency was always there. The update just turned it up loud enough to see. And it isn’t one app’s quirk: every assistant, Claude included, is trained on human approval, so every one of them leans agreeable.

THE SYCOPHANCY PROBLEM

 **Why it happens**

Models are tuned on human ratings. Friendly and validating wins those ratings. Honest pushback loses them. The model isn’t lying to you. It learned what people give a thumbs-up.

 **What to watch for**

Enthusiasm with no specifics. Instant agreement with your opinion. Caving the moment you push back. If the feedback never touches your actual words, it isn’t feedback.

WHAT THE INDUSTRY IS DOING

The labs aren’t ignoring this. Since the rollback, sycophancy is something they measure: new models get tested for it before they ship, the way they get tested for accuracy.

The fixes run in two directions. On the training side, better rating systems that stop rewarding empty agreement. On the product side, assistants tuned to push back more, with settings that let you ask for a blunter style.

None of it is solved. The pull comes from training on human approval, and that isn’t going away. So the last guardrail is you, and it’s mostly about how you ask:

MAKE IT A COACH

A hype man feels better. A coach makes you better. Three prompts turn the hype man into a coach:

1 ? Ask, don’t tell

Researchers measured it: open with “I think this plan is good” and the model tends to agree with you. Ask the neutral question instead and the flattery mostly disappears. The more conviction in your prompt, the harder it caves. Hide your hand.

“Here’s my plan. Evaluate it before I tell you what I think.”

What you’ll get: Same plan, same model. Tell it your opinion first and it builds your case. Ask cold and it judges.

2 🔍 Ask for the gaps

Models default to encouraging. This tells it the job is different: you want the gaps, not the highlights.

“Don’t tell me what’s good yet. What’s weak, what’s missing, and what would someone who disagrees say?”

What you’ll get: Instead of “This is a great start!”, you’ll get “Your thesis is vague, paragraph 3 contradicts paragraph 1, and your conclusion introduces a new argument instead of wrapping up.”

3 📋 Grade against a rubric

A rubric forces the model to judge against a standard instead of reacting to your effort. Especially useful for school work, where the rubric usually exists anyway.

“Grade this against the following rubric. For each category, give me a score, quote the part of my work that supports the score, and explain what would move it up one level. [paste rubric here]”

What you’ll get: Instead of “this is a strong essay,” you’ll get “Thesis: 3/5. Your claim is clear but not specific. To reach 4/5, name what specifically about the Dream Gatsby exposes.”

4 ✂️ Make it argue the other side

If AI agrees with your position, flip it. If it can’t build a strong case against you, your position might be solid. If it can, you’ve found your blind spots.

“I think [your position]. Now argue against me as strongly as possible. Give me the 3 best counterarguments someone could make, and explain why a reasonable person might hold each one.”

What you’ll get: You’ll find the holes in your argument before your teacher or debate opponent does.

ONE MORE THING

You can accidentally bake the Flattery Trap into your settings. Personalization like “*always be supportive*” or “*validate my thinking*” feels good, but it makes the AI worse at the criticism you actually need.

Try instead: “*Be supportive, but don’t flatter me. Challenge weak reasoning and point out what needs work.*”

 **Don’t take praise as a verdict.** AI is still a strong editor. It can find the vague thesis, the missing counterargument, the edge case you skipped. It just won’t volunteer any of that while it’s being nice. Warmth with no specifics is your cue: ask what’s weak, what’s missing, what a skeptic would say. Then take the praise, if it survives.

Friendly isn’t the same as right.

The warmer it sounds, the harder you check.