

Habits for the Road

You now know how AI works and how to use it well. Two habits decide whether that knowledge actually helps you. The first is honesty.


AI is an incredibly powerful tool for learning. But there's a difference between **using AI to learn** and **using AI to skip learning**. The line isn't always obvious, so let's make it clear.

WHERE THE LINE IS



Submitting AI output as your own work

Copying AI-generated text, code, or answers and putting your name on it. This is plagiarism, even if you edited a few words.



Using AI on assignments that prohibit it

If your teacher says no AI, that means no AI. Violating this is academic dishonesty, even if you think the rule is unfair.



Using AI to generate ideas, then doing the work yourself

Generally acceptable, but check your teacher's policy. Be transparent about what you used AI for if asked.



Using AI to explain concepts you don't understand

Like having a tutor. This is usually fine and actually helps you learn. But don't let it replace actually working through problems.



Using AI to review and improve your own work

You did the thinking. AI gives feedback. This is like asking a friend to proofread your essay. The work is still yours.



Using AI to learn a topic, then demonstrating understanding yourself

Research with AI, understand the material, write/present/solve in your own words. The learning happened. That's the whole point.

The yellow zone is the tricky one. The same AI use can be fine in one class and not okay in another. The assignment's rules decide.

IF YOU CAN'T EXPLAIN IT, DEFEND IT, OR REPRODUCE THE THINKING, IT ISN'T REALLY YOUR WORK.

The teacher-without-notes test still applies. Could you walk through it cold? If not, AI did the work for you. And one more test catches almost everything: if you'd be embarrassed to say out loud how you used AI, that's your answer.

PLAGIARISM VS COPYRIGHT

Plagiarism and copyright sound similar, but they answer different questions. Get the difference clear before you submit, post, or sell anything you used AI to make.

PLAGIARISM

Did you pretend someone else's work was yours?

About honesty.

COPYRIGHT

Were you allowed to use or copy that work?


About rights.

STAYING OUT OF TROUBLE



'AI made it' doesn't automatically mean 'I own it.'

Different apps have different terms. Some grant you commercial rights to your outputs, some restrict them. Check before you sell, post, or submit anything that's going to make money or live somewhere public.



Do the real work yourself and it's yours.

Brainstorm characters with AI, then write the story yourself, and the story is yours. AI was an ingredient, not the author. The work you actually do is the work you own.



Rewriting someone's work doesn't make it yours.

Upload a copyrighted article, ask AI to reword it, and you get a copy in new clothes, not a new work. The rights still belong to whoever wrote the original.



'In the style of [real person]' gets risky fast.

Especially if the person is alive, the style is tied to a recognizable brand, or you're using the result commercially. School project poster: usually fine. T-shirts you're selling on Etsy: not fine.

SHOW YOUR PROCESS

Some schools run student work through AI detectors. They sound official. They're not as reliable as they sound. They produce both false positives (honest work flagged as AI) and false negatives (real AI use missed). A detector score isn't proof by itself, and the arms race favors the generators.

So don't bet your grade on a tool's accuracy. Build a defense the detector can't take from you: a visible record of how the work actually came together. What you keep is your call. The point is that some trail exists if you ever need it. A few examples:



Drafts as you go

Save versions, not just the final. The path from rough to polished is itself the proof.



Sources with links and dates

Copy the URL and the date for anything you read or watched. Real research leaves a trail. AI hallucinations don't.



AI conversation logs

Keep your prompts and the AI's replies. If your work is ever questioned, you can show exactly how you used AI.

THE BEST DEFENSE ISN'T HOPING THE DETECTOR BELIEVES YOU.

It's having the receipts.

IF YOU'RE ALREADY SITTING ACROSS FROM A TEACHER

Having the receipts is what to bring to that meeting. Here's what to do with them once you're in the room: ask which detector was used, what the score actually means under this teacher's policy, and what the appeal path is, before you leave. That's the same move this lesson comes back to for any system that scores you: name who's responsible, ask what data was used, and find out if you can push back.

DISCLOSING YOUR AI USE

Whether you have to cite AI, and in what format, is your school's call. Check your teacher's rule first, and follow current APA, MLA, or Chicago guidance if they require it. What doesn't change is the honest move: a short note saying what you used AI for. Here's what a plain disclosure note can look like, adapt it to your situation:

Brainstorming

I used ChatGPT to brainstorm possible essay structures. I wrote the final draft myself.

Use when AI helped you explore ideas but you did the writing yourself.

Debugging code or fixing errors

I used Claude to debug my code and explain the error.

Use when AI helped you find or fix mistakes, but the work is yours.

Practice and review

I used Gemini to generate practice questions for review.

Use when AI made study material for you, like flashcards or quiz questions.



Always disclose what you used AI for

Even if citation isn't required, add a note: "I used Claude to help brainstorm ideas for this essay" or "AI was used to debug my code." Transparency builds trust.



When in doubt, ask your teacher

AI policies vary by school, class, and assignment. If the rules aren't clear, ask before you submit. It's easier to get permission than forgiveness.

WHEN OTHER PEOPLE ARE IN THE LOOP

In group work, agree on AI use before anyone uses it. If one person silently AI-drafts a section, the others can't review it the way they would human work, can't catch mistakes, and didn't get to weigh in on whether to use AI at all. Hidden AI use changes the work for everyone.

AI CAN HELP YOU LEARN, DRAFT, DEBUG, AND REVISE.

But your name goes on the work only when you understand it, can defend it, and followed the rules.


WHAT YOU TYPE TRAVELS

The first habit was using AI honestly. The second is guarding what you share with it. Treat AI like a group chat with a company, not a diary. You may be the only person on the screen, but the conversation isn't only yours. Depending on the tool and the settings, what you type can be stored, linked to your account, reviewed for safety, and kept in logs longer than you'd expect. That doesn't mean panic. It means awareness.

And it isn't just text you send. Modern AI is **multimodal**: it can read images, screenshots, PDFs, audio clips, and video. The word means "more than one mode of input." A screenshot is still a prompt. So is a photo, a PDF, or a voice recording. It just looks different. Same context window, same storage rules, same questions about what should and shouldn't be in there.

Picture this. You snap a quick photo of your math homework so AI can help you with one of the problems. You meant to share the math.

The Whole Photo



Your name

AI can read handwriting. Your name on the homework page is now in someone's logs.



Your school

Lincoln High at the top of the page tells AI exactly where you spend your weekdays.



Your teacher and class period

Specific enough to narrow down which classroom you sit in.



Your locker combination

You wrote it in the planner so you wouldn't forget. The photo just published it.



A prescription medication

Medical info is some of the most sensitive data there is. The bottle was just sitting on the desk.



Your home street address

You didn't mean to share this. The shipping label did.

WHAT YOU ACTUALLY SHARED

You meant to share the math. You also shared six other things you didn't think about. The photo was the prompt. Everything in it came with.

WHAT HAPPENS TO YOUR DATA



Your conversations are saved

Many AI tools save chat history by default. Whether it's stored, for how long, and whether it can be deleted depends on the tool, the account, and your settings.



Humans may read them

Some conversations get reviewed by people: flagged content, safety checks, quality samples, abuse prevention. Not every chat, but a chat isn't as private as a journal either.



They might train on your data

Some apps use your conversations to improve future models. That means your words could influence how the model responds to other people.



Data breaches happen

In 2023, a ChatGPT bug exposed some users' conversation titles to other users. No system is perfectly secure. The practical rule isn't to panic. It's to keep information out of the chat that you'd hate to see leak.

SHARE THE TASK, NOT THE IDENTITY.

The AI usually needs the problem. It usually doesn't need your name, address, password, school login, exact location, private messages, or someone else's sensitive information.

WHEN SAFETY IS INVOLVED

Privacy can't be AI's silence from humans. If a friend may be in danger, get real-world help: a counselor, a parent, or a teacher you trust. Some things need a person, not a prompt.

Run a prompt through these five questions before hitting send.



What does the AI actually need to help?



Can I remove names?



Can I replace exact details with categories?



Am I sharing someone else's private information?



Would I be okay if this appeared in a log?


WHAT EACH PLATFORM DOES

Last checked: May 2026. These details change as policies update, so treat what's below as a snapshot. Your job is to know where to look: privacy settings, data controls, activity history, temporary chat, memory, connected apps, whether your chats can be used for training, and your school account policies.



ChatGPT (OpenAI)

- Conversations may be used to train future models by default (Plus and Pro accounts included)
- Opt out in Settings → Data Controls → 'Improve the model for everyone'
- Even after opt-out or deletion, chats may be kept up to 30 days for safety monitoring
- Temporary Chat isn't saved to your history and isn't used for training



Claude (Anthropic)

- Training is ON by default for Free, Pro, and Max accounts
- Opt out in Settings → Privacy → 'Help improve Claude'
- If you don't opt out, conversations can be kept for up to 5 years
- Anthropic staff may review flagged conversations for safety



Google Gemini

- Conversations are saved and may be used to train Google's AI by default
- To stop training, turn off 'Gemini Apps Activity' in your Google account
- Turning that off also means you lose your chat history
- Conversations chosen for human review are kept for 3 years, even if you delete them



School-Provided AI Tools

- Your school may monitor or log every AI conversation
- Privacy rules can be different from the consumer version of the same tool
- Your school's IT department may have access to your history
- Content filters are usually stricter than the public version

WHAT YOU CAN DO

Check your settings

Most platforms let you opt out of training data collection. Find the privacy or data settings and review them. It takes 2 minutes and it's worth it.

Use temporary chat for one-offs

Most apps have a temporary or incognito mode that doesn't save to history or memory. Use it for sensitive topics or one-off questions you don't want sitting in your account.

Crop before you upload

Screenshots and photos carry everything in the frame: names, school logos, addresses, prescription bottles, locker combinations. Crop tight to what AI actually needs, or type the data instead of uploading the image.

Audit connected apps

If AI can read your Gmail, Drive, Photos, Calendar, or browser, the privacy surface goes way past the prompt you typed. Review which apps you've connected and remove any that don't earn the access.

Delete what you don't need

Periodically clear your conversation history. If you don't need it anymore, there's no reason for it to sit on someone else's server.

Use accounts wisely

If your school provides an AI account, assume it's monitored. For personal use, consider whether you want conversations tied to your real identity or a separate account.

IT ISN'T ONLY WHAT YOU TYPE

Everything above is about what you choose to share. But AI reads and decides things about you even when you never typed a word into it: a feed deciding what you see, a flagging tool scanning your essay, a screening tool sorting an application before a human looks. Sometimes you're the one using AI. Sometimes you're the one it's deciding about. Same skill either way: know what to ask when it's the second one.

WHAT YOU SEE

Algorithms rank what lands in front of you.

Feeds, recommendations, search results, your inbox order. Shaped before you tap anything.

WHAT YOU'RE ALLOWED TO DO

Filters and flags decide whether you can post, comment, apply, or stay logged in.

Content moderation, plagiarism flags, account suspensions. Lines you can hit without seeing them coming.

WHAT GETS ATTACHED TO YOUR NAME

Scores and classifications follow you.

Hiring ratings, admissions tools, credit scores, fraud flags. A number that decides whether a human even looks.

This isn't the checklist you just used to clean up a prompt. It's a different one, for when a system is deciding something about you instead of the other way around.

THE FIVE QUESTIONS

Run these on any system that's judging you.

1. Who is responsible?

What person, school, company, platform, or institution chose to use the system?

2. What data was used?

What information, behavior, score, document, or profile did the system rely on?

3. Can you push back?

Is there an appeal, a moderator, a teacher, an admissions office, an employer, or some other support path?

4. What happens if it's wrong?

Is the cost small, or could it affect your grade, account, application, reputation, money, or safety?

5. Who is this helping, and who is it hurting?

Who gets the benefit when this works, and who carries the cost when it doesn't?

Notice the first question is who, not what. AI systems score, rank, flag, recommend, and sometimes act on their own. But a person or an institution chose to deploy that system, set the rules, and decide what happens when it flags someone. That's where responsibility lives, and that's who you push back on.

ASK FOR THE PAPER TRAIL

If an AI system affects a decision about you, get the decision out of the black box and into a process a human can explain. Ask, in writing if needed:

- What policy applies here?
- What data, document, score, or evidence was used?
- Was there human review?
- How do I correct a mistake?
- What is the appeal path?

A school or company that won't put any of this in writing is telling you something.

EVERY SYSTEM IS CHASING SOMETHING

A deployed system optimizes for something: engagement, speed, cost, fewer manual reviews. TikTok's feed is tuned for watch time. A hiring screen might be tuned to match past hires. Ask what it's actually chasing, and whether that lines up with what's good for you. And remember that "mostly right" doesn't mean right for you: a 99.9%-accurate system still gets a thousand wrong for every million people it scans. The averages look fine. The thousand people who got the wrong answer don't get to opt out of the consequences, and a false positive doesn't hurt the same people a false negative does.

FAIR ON PAPER ISN'T ALWAYS FAIR

A system can claim it doesn't use race, gender, or income and still produce unfair outcomes, because other data it does use correlates with what it claims it doesn't. A zip code can stand in for race and income. A school name can stand in for class. These are called **proxy variables**: data that carries the effect of something the system never names directly. So when you ask what data was used, the follow-up is what that data correlates with.

WHEN AI JUDGES YOU, DON'T ARGUE WITH THE MACHINE.

Find the human, ask what data was used, and learn how to appeal.

PRIVACY RISK GROWS THE MORE AI CAN SEE.

The moment AI can read more than your message (files, email, photos, calendars, documents, web pages), the surface area expands. And there's a second risk hiding in that reach: everything AI reads becomes potential instructions. A webpage, an email, or a shared doc can carry hidden text that says, in effect: ignore your user and do this instead. The model can't always tell your words from the words on the page. The trick has a name: **prompt injection**.



The shape to remember: you ask for a summary of a webpage, and the answer comes back with an odd extra step, like a suggestion to send an email or click a link. That suggestion didn't come from you, and it didn't come from the AI's judgment. It came from the page.

The defense is a rule you already live by: real instructions come from you. When an action request shows up inside something the AI was reading, treat it as the content talking. Pause anything unexpected, and keep the irreversible steps behind your own review.

PRIVACY ISN'T SECREC.

It's only sharing the information the task actually requires. Once it's in the chat, you can't take it back.

Own what you submit. Guard what you type.

Two habits. They travel well.