


How an LLM Works

You've met the two kinds of AI and named the generative one: a Large Language Model, or LLM. So how does an LLM actually turn your words into an answer?

When ChatGPT or Claude write a sentence, they're running math to predict the likely next words. They aren't looking up what your words mean; they're working out which words tend to follow which.

How do they do it? In two phases. First the model learns, once, by soaking up patterns from mountains of text. Then, every time you chat, it uses those patterns to build your answer one word at a time.



Four ideas carry it. Below, we follow one example, peanut butter, through all four.

01 🌱 Training LEARN ONCE

The model **teaches itself**: guess the next word, check, and nudge its numbers toward the right word.



Repeat that loop billions of times, and that's training.

02 🍀 Patterns LEARN ONCE


So what is it actually learning? **Patterns**. Here's one you picked up as a child. Which word comes next?



AI didn't memorize these. It absorbed the pattern by running through billions of examples.

03 🎲 Probability EVERY WORD


AI doesn't make one guess. It scores **every** possible next word: a ranked list with a probability on each, and those numbers shift with the surrounding text.




Add the word banana, and jelly's probability drops from 41% to 2%. The word banana changed everything.

04 🍷 Prediction EVERY WORD


Probability handled one word. But your answer is hundreds of words long, so the model just repeats the move. Your phone does this when you write a text: it suggests a word, you tap it, it suggests the next. AI works the same way, with no fixed plan for where the sentence will end up. One word, look again, the next.



A full paragraph runs this loop many times, fast enough to look like thought.



That clears up three big myths:



Not magic. Not a person. Not a truth machine.

It's math working out probabilities.