

Judging AI's Future

Some people think AI will cure cancer, eliminate poverty, and make everything better for everyone. Others think it will steal every job, seize control of infrastructure, and end civilization as we know it. Both camps are completely confident. Both can't be right. Here's how to figure out who's actually worth listening to, and how to make that call yourself from now on.

The terms you'll hear in the debate

Everything you've used so far in this course is narrow AI, or ANI: software built for one specific job. Read enough about where AI is headed and you'll hit two more acronyms, both about getting past narrow.

AGI (Artificial General Intelligence): a hypothetical AI that can do any intellectual task a human can. Some researchers think it's a few years away. Others think it's decades. Others think the term is meaningless.


ASI (Artificial Super Intelligence): a hypothetical AI smarter than any human, at everything. Even more contested.

Neither has clearly emerged yet, and the definitions are still contested. Both shape how people argue about AI today. The vocabulary matters, but it isn't the real skill. The real skill is judging the claims people make in either direction. The rest of this lesson teaches three tests for doing exactly that.

History doesn't say ignore big predictions. It says ask what gets ignored between possibility and reality.

The pattern: predictions fail on friction.

Most predictions fail because people confuse what is technically possible with what can actually spread, scale, and stick.



Technical possibility is not social inevitability.

HOW TO JUDGE THE NEXT CLAIM

Run every claim through three tests.

When you hear a prediction about AI (scary or exciting), do not just ask whether it sounds possible. Ask what would need to be true, and what still gets in the way.

1 Incentives

Who benefits if I believe this?

A politician says AI will destroy your future. A tech CEO says AI will save the world. A news site says AI is coming for your job. Ask: are they trying to get your vote, your investment, or your click? The claim might still be true, but you should know what's driving it.

2 Mechanism & Friction

Does this fit how AI actually works?

'AI means you don't need school anymore.' For that to be true, every AI answer would have to be correct, colleges and employers would have to stop valuing degrees, and you'd have to learn complex skills with no structure, feedback, or deadlines. None of that is happening.

3 Evidence

Is this based on evidence, or just extrapolation?

'AI beat a pro gamer at chess' is evidence. 'AI will dominate all esports' is extrapolation. One controlled game is not the same as real-time team strategy in Valorant or Rocket League. A demo proves something is possible. It doesn't prove it will spread, scale, or dominate soon.

A claim being overhyped doesn't make it fake, and a claim being real doesn't make it soon. The three tests are how you tell which.

AI IS A BLACK BOX

There is a term software engineers use for AI: **the black box**. With a regular calculator, if you type 2+2 and it says 5, a programmer can open the code, find the broken line, and fix it. The path from input to output is traceable.

AI is different. When you give ChatGPT a prompt and it writes you a poem, nobody (not even the people who built it) can point to a specific spot inside the model and say "that's where the rhyme came from." You can study its behavior from outside, but you can't crack it open and read the logic like normal software.

So what IS inside? Learned numbers. Billions of them, adjusted during training until they helped the model predict better.



That opacity is exactly why future claims are hard to check. If someone says "AI will definitely do X soon," the technical-sounding explanation is not the evidence. The evidence is what shows up outside the model: real-world performance, repeated results, deployment at scale, known failure rates, accountable owners.



The rules are learned, not written.

A person didn't write every possible answer by hand. The model learned patterns from training data.



The inside isn't a repair manual.

Even if you could read billions of numbers inside the model, they wouldn't look like clear instructions.



Its explanation may not be the real cause.

If you ask why it said something, the answer can sound reasonable but may not actually trace the internal cause.



Closed tools hide even more.

With ChatGPT, Claude, and Gemini, you don't see the full model, training data, system instructions, or safety layers.

A confident explanation isn't the same as evidence. Outside-the-model evidence is what tells you whether a future claim is real.

Predictions are cheap. Reality has friction.

Judge the claim by what it leaves out.