

# Opener

THE TRAPS AHEAD

The false fact sounds sure.  
The flattery feels good.  
The fake looks real.  
— *every trap looks fine from the inside.*

You just spent a whole section under the hood. Now comes the warning that goes with it: the same machinery that makes AI powerful also fails in specific, predictable ways.

And here's the catch. When it fails, nothing looks broken. A made-up fact reads exactly like a real one.

Think about a rip current in the ocean. It doesn't look dangerous. It looks like the best spot on the beach: a calm, flat channel between the breaking waves.

Swim into that calm, and you'll be pulled far from shore. The skill isn't out-swimming a rip current. It's recognizing one before you're in it. Lifeguards spot them all day, not because they're strong swimmers, but because they know the shape a rip current makes in the water.

AI is the same. It has its rip currents, and recognizing them is the real skill. This section teaches you to read the water.

And you have a head start most people never get. A rip current isn't the ocean malfunctioning. It's just waves doing what waves do. Every AI trap ahead works the same way.



IN THIS SECTION

Avoid Traps

TRAPS IN THE ANSWER

- 1
- First, the ways an answer goes wrong all on its own: facts invented with total confidence, bias and stale knowledge soaked up from the training data, and summaries of documents the model never really read.

TRAPS IN YOU

- 2
- Then, the traps that work on you instead of the answer. Helpful, agreeable, engaging: the qualities that make AI easy to use also make it easy to fall for.

TRAPS FROM THE WORLD

- 3
- Finally, the trap that comes looking for you: other people's AI, putting fakes in front of you so convincing that seeing is no longer proof.

When it fails, nothing looks broken.

*This section teaches you to read the water.*