

Opener

THE TRAPS AHEAD

Spot the trap.
Name the trap.
Avoid the trap.

— *every one looks fine from the inside.*

You just spent a whole section under the hood. Now comes the warning that goes with it: the same machinery that makes AI useful also fails in specific, predictable ways. And when it fails, nothing looks broken. A made-up fact reads exactly like a real one. Flattery feels exactly like insight. The trap that catches you won't announce itself. It will look fine.

Think about a rip current. It doesn't look dangerous. It looks like the best spot on the beach: a calm, flat channel between the breaking waves. That calm is the trap. Lifeguards are strong swimmers, sure. But strength isn't what spots a rip. Recognition is. They know the shape a rip makes in the water, so the calm patch never fools them. This section teaches you to read the water.



And you have a head start most people never get. A rip current isn't the ocean malfunctioning. It's just waves doing what waves do. Every AI trap ahead works the same way. Hallucination is prediction doing what prediction does. A skewed answer is data doing what data does. Flattery is a model trained to please, pleasing.

IN THIS SECTION

Avoid Traps

TRAPS IN THE ANSWER

- 1 First, the ways an answer goes wrong all on its own: facts invented with total confidence, bias and stale knowledge soaked up from the training data, and summaries of documents the model never really read.

TRAPS IN YOU

- 2 Then, the traps that work on you instead of the answer. A tool built to be helpful, agreeable, and engaging can feel like a mind, cheer on your worst ideas, hold your attention longer than you meant to give it, and stand in for support it can't actually provide.

TRAPS FROM THE WORLD

- 3 Finally, the trap that comes looking for you: other people's AI, putting fakes in front of you so convincing that seeing is no longer proof. Named traps are avoidable traps.