

Training Bias

Sometimes AI doesn’t invent anything. Every fact in the answer is real, but the picture is still distorted.

One big cause is training bias. The model learns whatever patterns are in the data it saw. If the data is uneven, incomplete, or full of shortcuts, the model can treat a narrow slice of reality as the default.

Imagine the only driving advice online came from professional race car drivers. If you asked, “What’s the best way to drive?” the model might tell you to drive 180 miles per hour on a closed track.

That answer wouldn’t come from evil intent. It would come from a narrow slice of data treated like the whole picture.

COWS ON THE BEACH

Computer-vision researchers hit a famous version of this. They had a model that could spot cows in photos with high accuracy. Then they showed it a cow on a beach, and it went blank. Same animal, same shape, same spots.

The model was training on photos that almost all showed cows on green pasture. So, the model learned the wrong pattern: green grass means cow. It never learned the animal. It learned the background.



HOW BIAS GETS IN

Not one mechanism. Four overlapping ones, all rooted in the data the model was trained on.

**Overrepresentation**

Some cases appear too often in the data. The model treats them as the default.

**Underrepresentation**

Some cases appear too rarely. The model performs worse when those cases come up.

**Shortcuts**

A wrong clue happens to work in training. The model learns the clue instead of the real concept.

**Biased answer key**

Even balanced data can be graded wrong. If biased human decisions were the answer key, the model learns that bias as correct.

There’s no fact-check for this trap: every fact in a distorted answer can check out. That’s why your Evaluate the Results dig list asks who’s missing. Fact-checks cover what’s on the page; bias hides in what never made the page.

STALE INFORMATION

One more data trap, and this one isn’t bias or hallucination. Training stopped on a cut-off date, and the model has been answering from that snapshot ever since. Ask about anything newer, and it answers as if time never passed. It isn’t inventing. It’s remembering, and the memory is out of date.

The fix: for anything that can change, check the date. The model often runs a live search on its own, but not every time, so when it matters, ask for the search yourself.

The model learned our shortcuts.

Including the ones we’d rather it didn’t.