

Tune the Model

ChatGPT, Claude, and Gemini are all great choices for your go-to AI app. Now comes something normal apps never ask of you. Spotify doesn't ask which Spotify you want. AI apps do: you get choices about which model does the work, how fast it answers, and how hard it thinks.

Think about buying a new car. Your first decision is which dealership to visit: Ford, Chevy, Dodge, or somewhere else. When you arrive, the next choice is which model you want to buy.

Next come the options: engine size, wheel package, trim. And you'll pick those based on how you plan to use the car.

AI works the same way. The app was the dealership, and you've already picked it. The choices waiting inside are the model, the speed, and the thinking.

YOUR CHOICES

The choices live inside whichever app you picked. What makes them confusing is the names: each app names its models differently, and with every app improving this fast, the names keep changing. So don't memorize names. Learn the three choices underneath them. Those don't change.

CHOICE 1 Model	CHOICE 2 Speed	CHOICE 3 Thinking
WHAT IT IS A family of models, not one brain. The picker lists them by name, newest near the top. The labels carry the signal: one is marked most capable, one is the everyday default.	WHAT IT IS Trimmed versions of the brain, built to answer fast. Short names that sound quick: mini, flash, lite. Some apps list them as their own models, others fold them into one menu.	WHAT IT IS A switch that lets the model work longer before it answers. Named things like extended thinking, reasoning, or think longer. Same brain, more time.
WHY IT EXISTS Bigger models reason deeper and cost more to run. Smaller ones are cheap enough to answer everything else. No company can give everyone the biggest brain for every question, so the picker exists.	WHY IT EXISTS A quick question doesn't need the deepest brain. Fast versions trade some depth for speed, and on most quick asks you'll never notice what they traded.	WHY IT EXISTS Hard problems need more than one pass. Thinking time lets the model try approaches, check its own work, and throw out weak drafts before you ever see one.
WHEN TO USE IT The default handles most days. Switch up when the work is hard, long, or graded.	WHEN TO USE IT Quick lookups, simple rewrites, anything you'd be fine asking twice. When the answer matters more than the wait, move up.	WHEN TO USE IT Math, code, planning, anything with steps that build on each other. Skip it for quick facts: you'll wait longer for the same answer.

One thing worth knowing about that first choice: the models in a family aren't one brain dialed up and down. They're different LLMs. The big one has more layers, more dimensions, and more training behind it. The small ones were built lighter from the start.

All three choices live in the same corner of the screen. Look for the model name at the top of the chat or next to the message box, and tap it: the whole family drops down, usually with the thinking switch close by. If there's no picker at all, the app is choosing for you.

THE DEFAULT IS A CHOICE SOMEONE ELSE MADE.

Most people never open the picker and never learn what's in there. You just did. Open it and choose on purpose when the problem deserves it.

One more thing if you're on a free plan. The top tier has the tightest limits, and when you hit one, the app may quietly switch you to a smaller model or make you wait. If answers suddenly feel dumber mid-chat, that's usually what happened: check the picker. And limits usually burn by how much text is in the conversation, not how many messages, so long pastes cost more. Start fresh chats when you can.

THINKING, SIDE BY SIDE

Thinking is the choice people skip most, so here are the two settings next to each other. And it isn't the same as picking a model, even though both say 'turn it up for hard problems.' A bigger model is more brain. More thinking is more time from the brain you already picked. Rule of thumb: pick the model that fits the job first. If the answer still feels shallow, turn the thinking up before you switch models.

 Quick answer	 Deeper thinking
How it works: Pattern matches from training, answers immediately. Best for: - You need a quick answer - You're brainstorming lots of options - The task is simple formatting, rewriting, or summarizing - You're asking for a first draft, not final judgment - Speed matters more than depth	How it works: Reasons step by step before responding, checks its own work. Best for: - The problem has multiple steps - The answer needs logic, math, code, or careful comparison - You're analyzing an argument, plan, essay, or decision - The first answer felt shallow and the issue is reasoning - Accuracy matters more than speed - The decision has real consequences

MORE THINKING ISN'T A TINY PERSON INSIDE THE COMPUTER RECONSIDERING.

It's an extra round of reasoning and self-checking the model runs before answering, which catches more on hard problems. But it's not a truth machine. If the model is missing information, using old information, or relying on a bad source, thinking longer won't fix that. For truth, you still need to evaluate the results.

THE DIAL YOU WON'T SEE

Three choices are yours. This one isn't. Back in One More Thing you met the weighted drawing: every token holds tickets equal to its probability, and the model draws one. **Temperature** is the dial on that drawing. Turn it down and the top pick takes nearly every ticket, so the model plays it safe and hands back almost the same answer every time. Turn it up and the tickets spread far down the list, so a long shot like Biscuit, which never wins at low, starts showing up. More range, more surprise, and sometimes more nonsense.

You won't find it in ChatGPT, Claude, or Gemini. The apps set it for you, and the dial itself lives in the tools developers use to build on these models. You can still ask for the effect in words. Ask for "your single best answer" and you pull toward the safe end. Ask for "ten very different options" and you pull toward the wild one. That's the one dial you can't reach. The rest of the machine you can.

App. Model. Speed. Thinking.			
One decision you made back in Work With AI, and three you make inside it every day.			
1 App Where you work. Your home base, plus another app when the job clearly fits its strength.	2 Model Which brain does the work: the most capable one, or the everyday default.	3 Speed How fast you need it back. Trimmed versions answer quick, and on simple asks you won't miss what they left out.	4 Thinking How long it works before it commits: off for quick facts, on for problems with steps.

A bad answer doesn't always mean 'try harder.' Sometimes you picked the wrong app. Sometimes the model is too small. Sometimes you're on the fast version. Sometimes it just needs more thinking. Knowing the four decisions means you can fix the right one.

The app was just the dealership.

The real choices are inside.