

Tune the Model

You’ve already learned to choose your app. Now comes something normal apps never ask of you. Spotify doesn’t ask which Spotify you want. AI apps do: inside each one is a family of models, and a picker to choose between them.


Think of a hockey team changing its lines. The first line has the most skill, but you can’t play them every shift. Their minutes are limited. The third line is fast and steady, and it takes the everyday shifts. The coach’s job is matching the line to the moment. Inside an AI app, you’re the coach.

And the lines are real. The models in a family aren’t the same brain dialed up or down. They’re actually different LLMs. The big one has more layers, more dimensions, and was trained on more data. The smaller ones are built lighter from the start.

Why build a family at all? Because every answer you get runs on a server doing heavy math, and that math costs real money, time, and energy. So every model is a trade between three things: how smart it is, how fast it answers, and what it costs to run. No model maxes all three. The family exists so you can pick where to sit on that scale.


THE MODEL FAMILY

Every model in a family sits somewhere on one scale. Nothing maxes both ends: picking a model is picking where to sit.



Notice we didn’t print any model names. That’s on purpose: the companies rename and replace models a few times a year as they compete, and a name alone won’t tell you its size anyway.

What you need is a way to figure it out in any app, in any year. Three moves:



THE DEFAULT IS A CHOICE SOMEONE ELSE MADE.

Most people never open the picker and never know what they’re missing. You know the tiers now. Match the line to the moment, and send out the first line when the problem deserves it.


One more thing if you’re on a free plan. The top tier has the tightest limits, and when you hit one, the app may quietly switch you to a smaller model or make you wait. If answers suddenly feel dumber mid-chat, that’s usually what happened: check the picker. And limits usually burn by how much text is in the conversation, not how many messages, so long pastes cost more. Start fresh chats when you can.

You’ve chosen your app, and you just chose your model. Now there’s one more choice: how much effort should the model spend on this answer?

Every major AI lets you set how much effort it spends before answering. Turn it down and the model answers fast from what it already knows. Turn it up and it reasons through the problem step by step before committing.

Effort changes how hard the model thinks before it commits, not how the answer sounds. The wording has its own dial, and it’s coming up.

And it’s not the same as the model choice you just made, even though both say ‘turn it up for hard problems.’ A bigger model is more brain. More effort is more thinking time from the brain you already picked. Rule of thumb: pick the model that fits the job first. If the answer still feels shallow, turn up the effort before you switch models.



MORE EFFORT ISN’T A TINY PERSON INSIDE THE COMPUTER RECONSIDERING.

It’s an extra round of reasoning and self-checking the model runs before answering, which catches more on hard problems. But it’s not a truth machine. If the model is missing information, using old information, or relying on a bad source, thinking longer won’t fix that. For truth, you still need to evaluate the results.

WHERE TO FIND THE DIAL

Look near the model picker. Some apps give effort its own switch, with a label like ‘extended thinking’ or ‘reasoning.’ Others fold it into the picker itself: choosing a model labeled ‘thinking’ is choosing more effort. Read the labels the same way you read the picker, and when you can’t tell, a ten-second search settles it.


💡 Same dial, different labels in every app. Once you understand what it does, you can find it anywhere.

Some tasks need speed. Others need depth. And the dial turns both ways: deeper is slower, and on a simple ask it buys you nothing.

ONE MORE DIAL: TEMPERATURE

You’ll never touch this one. Back in One More Thing you met the weighted drawing: every token holds tickets equal to its probability, and the model draws one. **Temperature** is the dial on that drawing. Turn it down and the top pick takes nearly every ticket, so the model plays it safe and hands back almost the same answer every time. Turn it up and the tickets spread far down the list, so a long shot like Biscuit, which never wins at low, starts showing up. More range, more surprise, and sometimes more nonsense.

You won’t find it in ChatGPT, Claude, or Gemini. The apps set it for you, and the dial itself lives in the tools developers use to build on these models. You can still ask for the effect in words. Ask for “your single best answer” and you pull toward the safe end. Ask for “ten very different options” and you pull toward the wild one. That’s the one dial you can’t reach. The rest of the machine you can.



A bad answer doesn’t always mean ‘try harder.’ Sometimes you picked the wrong app. Sometimes the model is too small. Sometimes it just needs more effort. Knowing the three decisions means you can fix the right one.

Pick the model, set the effort.

Same dials, different labels everywhere.