

Check the Results

By now you’ve seen it yourself: AI makes mistakes. It states wrong facts in the same confident voice it states right ones. So what do you do with an AI answer, knowing it might be wrong?

One option is to ignore AI entirely. Nobody who’s made it this far in the course is doing that: AI is too powerful to leave on the shelf. The real move is in between trusting and dismissing. **You evaluate the results**, every time, with a process. Here it is.

FIRST, READ IT

This sounds too obvious to write down, and yet: weak AI users take AI’s output and share it without reading it. They paste it into the assignment, drop it in the group chat, hit send. Reading every word is the floor. Every other check in this lesson builds on it.

UNDERSTAND WHAT IT’S SAYING

You’ll never know whether an answer is right if you don’t understand what it says. If a paragraph uses a term you don’t know, or a step you can’t follow, that’s not a reason to skip it. It’s a reason to stop.

🔑 If something is unclear, ask. The fastest explainer of an AI answer is the same AI. “Explain the second paragraph in simpler terms” costs you one message, and now you can actually judge what it said.

VALIDATE IT AGAINST WHAT YOU KNOW

As you read, hold the answer up against your own knowledge. This is where everything you’ve ever learned pays off. Try it on these two AI statements. Both are false. Watch how differently they fail.

TWO FALSE STATEMENTS

YOU CATCH IT

“The American Civil War started in 1083, on the border of England and Argentina.”

Everything you’ve ever learned about the Civil War contradicts this: the century, the country, the geography. You don’t need to look anything up. Your own knowledge catches it instantly.

IT SLIPS PAST

“Maria Petronoski won 3 gold medals at the 1932 Olympics in Athens.”

Nothing about this sounds wrong. But the 1932 Games were in Los Angeles, not Athens, and Maria Petronoski doesn’t exist. Unless you happen to know Olympic history, “sounds fine to me” waves it right through.

The first statement shows why learning still matters in the AI era: a well-stocked head is your fastest fact-checker. The second shows its limit. When the topic is beyond what you know, “it sounds fine” isn’t validation, and you’ll need the deeper checks below. Hold that thought; the next lesson comes back to it.

DECIDE: DOES IT NEED MORE?

You’ve read it, understood it, and validated what you could. Before going deeper, ask two quick questions.

THE TWO-QUESTION GATE



1. What kind of task was this?

Factual work (research, real events, claims about the world) gets checked against the world. Generative work (first drafts, brainstorm) gets checked against what you asked for. Different task, different follow-up.



2. How much is riding on it?

Picking a movie for tonight? If it passed your read-through, you’re done. A college entrance essay, a health question, anything with your name on it? Keep going.

HOW MUCH TO CHECK

AI can help with almost any task. The real skill is knowing how much to check what it gives you. Some answers need a hard look. Others barely need one.

That’s what trust really comes down to here. Not believing the output automatically, but knowing what kind of follow-up it needs before you use it.

SOMETIMES THE ANSWER IS NO

Before the dial even matters, there’s a faster call. Some moments shouldn’t go to AI at all: the sharpest move is to leave the tool closed and handle it yourself, or hand it to a real person.

DON’T REACH FOR IT WHEN

THE RULE ALREADY SAYS NO

- The assignment or policy forbids it.
- You’d be using it to dodge the very skill you’re there to build.

IT INVOLVES PRIVATE INFORMATION

- You’d have to paste your own sensitive details.
- It means handing over someone else’s private information.

YOU CAN’T STAND BEHIND THE RESULT

- You can’t verify whether the output is true.
- It’s a high-stakes call that needs a qualified human.
- It would act under your name without your review.

IT’S YOURS, OR IT NEEDS A PERSON

- The task needs your own voice or judgment.
- Someone is in real danger and needs a person, now.

FOR EVERYTHING ELSE, SET THE DIAL

Everything else is fair game. From here, the question becomes how much follow-up the answer needs. Start by naming the **type of task** you are asking AI to do. Some task types are easy to judge on your own and need little follow-up. Other task types almost always need follow-up because the answer has to be true, current, safe, sourced, or genuinely yours.

How Much to Check

Can you judge the result yourself?

Yes



Good Fit

• Brainstorming

Generate options fast. You choose the best.

• Summarizing

Compress material you can check.

• First drafts

Get a starter version, then rewrite.

• Explaining concepts

Ask for a walkthrough at your level.

• Reformatting

Change the structure, not the meaning.

Best when you can judge whether the result makes sense.



Needs Follow-Up

• Citations

AI can invent sources that look real.

• Current events

Facts may have changed. Check a live source.

• Medical advice

Safety needs a real expert.

• Your Original Voice

AI can polish. You have to own it.

• Judgment calls

People, values, and consequences need you.

Answer needs truth, safety, sources, or your judgment.

Good Fit doesn’t mean ready to use. It means the task matches AI’s strengths and you can judge the result. The final call is still yours.

Now your turn. Below are prompts you might type. Make the call on each: is it a Good Fit, or does it Need Follow-Up?

🔑 The real question isn’t “Can AI do this?” It’s “What do I need to check before I use it?” That’s the move judgment makes.

GO DEEPER: FACTS THAT MATTER

A factual answer with real stakes needs evidence, not vibes. Three checks:

Check the citations. AI got its information from somewhere. Reply with: “Give me citations in your response.” The answer comes back with links to the pages behind it. Then, and this is the part people skip, click a link and read the page. Does it actually say what the AI claims?

Check that it’s current. As you’ll learn in the Training lesson, models have a training cut-off date: a point after which they know nothing. For anything recent (prices, standings, versions, news), ask the AI to run a web search and update its answer.

Know when it’s not AI’s call. Some questions deserve a professional, not a chatbot: medical symptoms, legal trouble, a mental-health crisis. AI can help you prepare questions for that person. It shouldn’t replace them.

Those three checks are the habit. The strategies below are the skill: how to verify a specific claim when it really counts.

TRUST, THEN CHECK

Is this true? Questioning an answer helps you spot a weak claim anywhere. AI hallucinates in specific ways that need their own toolkit. These six strategies will help you verify AI output before you use it.

You already learned to match the check to the task in How Much to Check: brainstorming, summarizing, and first drafts barely need a look, but facts, citations, and anything touching health, money, or another person do. This lesson is the other half. Once you know an answer needs checking, here is how you actually do it.

🔑 Confidence is not consistency. Two polished AI answers can disagree with each other, and both will sound certain. When you’ve asked twice and gotten different answers, don’t average them. The contradiction is the signal. Find the underlying claim and check it against something real.

The strategies below are tools. The judgment is which tool to use for which question.

STACK THEM, DON’T PICK ONE

In practice, you don’t run one strategy. You run several on the same answer. Cross-reference the specific claim. Check any citations. Watch for suspiciously specific numbers. Ask the model to flag what it’s uncertain about. The strategies above are the tools. Verifying is the workflow.

🔑 If you can’t find the source, assume it doesn’t exist. Confident wording is not evidence. Fancy formatting is not evidence. A real link to a real source you can read is evidence.

THE SIX CHECKS



Cross-reference the claim

Extract the fact. Search the fact. Don’t trust the sentence.



Verify citations and sources

If you can’t find the source after searching, it was likely fabricated. AI frequently invents plausible-sounding citations.



Don’t trust the source link

A sourced answer is better than an unsourced one, but a link doesn’t make a claim true.



Watch for hallucination triggers

The more specific the detail, the more skeptical you should be.



Ask the model to self-check

“Are you sure?” can surface uncertainty or catch mistakes, but it isn’t verification. Treat it as a second pass, then check the claim somewhere real.



Spot the ‘almost true’ trap

If part of an answer is verifiably true, don’t assume the rest is too.

GO DEEPER: DRAFTS THAT MATTER

A draft or brainstorm doesn’t need a citation trail. But it isn’t exempt from checking, because AI pads generative work with things you never said. For a draft that matters, ask three questions: does it do what I actually asked, does it fit the person who’ll read it, and did AI slip in any “facts” I never gave it? That last one is sneaky: ask for a game recap and it will happily invent the score of the third quarter.

🔑 True isn’t enough. The real question is whether the answer is good enough for what you’re about to do with it. A correct sentence in the wrong tone, at the wrong length, missing the point of your question, still fails.

THEN MAKE YOUR MOVE

Every evaluation ends in one of three moves.

THREE MOVES



Use it

It passed your checks. Most answers end here. Read, understood, validated, done.



Fix it

Something’s off and you can name it. Say exactly what’s wrong in a follow-up: “cut the stat I never gave you,” “make it shorter.” Being specific about what’s wrong beats starting over.



Walk away

Wrong tool, or stakes too high. Do it yourself, or take it to a person who actually knows: a teacher, a doctor, a pro.