

Check the Results

By now you've seen it yourself: AI makes mistakes. It states wrong facts in the same confident voice it states right ones. So what do you do with an AI answer, knowing it might be wrong?

One option is to ignore AI entirely. Nobody who's made it this far in the course is doing that: AI is too powerful to leave on the shelf. The real move is in between trusting and dismissing. **You evaluate the results**, every time, with a process. Here it is.

STEP 1 Read

This sounds too obvious to write down, but weak AI users take AI's output and share it without reading it. They paste it into the assignment, drop it in the group chat, hit send. Reading every word is the floor.

STEP 2 Understand

You'll never know whether an answer is right if you don't understand what it says. And the fastest explainer of an AI answer is the same AI: "Explain the second paragraph in simpler terms" costs you one message.

STEP 3 Validate

As you read, hold the answer up against your own knowledge. A well-stocked brain is your fastest fact-checker, and this is where everything you've ever learned pays off.

But watch its limit: "Maria Petronoski won 3 gold medals at the 1932 Olympics in Athens" sounds perfectly fine. The 1932 Games were in Los Angeles, and she doesn't exist. When a topic is beyond what you know, "sounds fine" isn't validation. That's what the deeper checks below are for.

STEP 4 Decide: does it need more?

Two quick questions before you go deeper.


- What kind of task was this? Factual work gets checked against the world. Drafts and brainstorms get checked against what you asked for.
- How much is riding on it? Picking a movie for tonight: you're done. A college essay, a health question, anything with your name on it: keep going.

And notice what Step 4 means: **sometimes read, understand, validate is all you do**. Low stakes and it held up? Use it and move on. The deeper checks are for answers that matter.

SET THE DIAL

How much follow-up does an answer need? Start by naming the **type of task** you asked AI to do. Some task types are easy to judge on your own and need little follow-up. Other task types almost always need follow-up because the answer has to be true, current, safe, sourced, or genuinely yours.

How Much to Check



Good Fit doesn't mean ready to use. It means the task matches AI's strengths and you can judge the result. The final call is still yours.

Now your turn. Below are prompts you might type. Make the call on each: is it a Good Fit, or does it Need Follow-Up?

 **The real question isn't "Can AI do this?"** It's "What do I need to check before I use it?" That's the move judgment makes.

GO DEEPER: FACTS THAT MATTER

A factual answer with real stakes needs evidence, not vibes. Three checks:

Check the citations. AI got its information from somewhere. Reply with: "Give me citations in your response." The answer comes back with links to the pages behind it. Then, and this is the part people skip, click a link and read the page. Does it actually say what the AI claims?

Check that it's current. As you'll learn in the Training lesson, models have a training cut-off date: a point after which they know nothing. For anything recent (prices, standings, versions, news), ask the AI to run a web search and update its answer.

Know when it's not AI's call. Some questions deserve a professional, not a chatbot: medical symptoms, legal trouble, a mental-health crisis. AI can help you prepare questions for that person. It shouldn't replace them.

Those three checks are the habit. The strategies below are the skill: how to verify a specific claim when it really counts.

TRUST, THEN CHECK

Is this true? Questioning an answer helps you spot a weak claim anywhere. AI hallucinates in specific ways that need their own toolkit. These six strategies will help you verify AI output before you use it.

You already learned to match the check to the task in How Much to Check: brainstorming, summarizing, and first drafts barely need a look, but facts, citations, and anything touching health, money, or another person do. This lesson is the other half. Once you know an answer needs checking, here is how you actually do it.

 **Confidence is not consistency.** Two polished AI answers can disagree with each other, and both will sound certain. When you've asked twice and gotten different answers, don't average them. The contradiction is the signal. Find the underlying claim and check it against something real.

The strategies below are tools. The judgment is which tool to use for which question.

STACK THEM, DON'T PICK ONE

In practice, you don't run one strategy. You run several on the same answer. Cross-reference the specific claim. Check any citations. Watch for suspiciously specific numbers. Ask the model to flag what it's uncertain about. The strategies above are the tools. Verifying is the workflow.

 **If you can't find the source, assume it doesn't exist.** Confident wording is not evidence. Fancy formatting is not evidence. A real link to a real source you can read is evidence.

THE SIX CHECKS

-  **Cross-reference the claim**
Extract the fact. Search the fact. Don't trust the sentence.
-  **Verify citations and sources**
If you can't find the source after searching, it was likely fabricated. AI frequently invents plausible-sounding citations.
-  **Don't trust the source link**
A sourced answer is better than an unsourced one, but a link doesn't make a claim true.
-  **Watch for hallucination triggers**
The more specific the detail, the more skeptical you should be.
-  **Ask the model to self-check**
"Are you sure?" can surface uncertainty or catch mistakes, but it isn't verification. Treat it as a second pass, then check the claim somewhere real.
-  **Spot the 'almost true' trap**
If part of an answer is verifiably true, don't assume the rest is too.

GO DEEPER: DRAFTS THAT MATTER

A draft or brainstorm doesn't need a citation trail. But it isn't exempt from checking, because AI pads generative work with things you never said. For a draft that matters, ask three questions: does it do what I actually asked, does it fit the person who'll read it, and did AI slip in any "facts" I never gave it? That last one is sneaky: ask for a game recap and it will happily invent the score of the third quarter.

 **True isn't enough.** The real question is whether the answer is good enough for what you're about to do with it. A correct sentence in the wrong tone, at the wrong length, missing the point of your question, still fails.

THEN MAKE YOUR MOVE

Every evaluation ends in one of three moves.

THREE MOVES



Use it

It passed your checks. Most answers end here. Read, understood, validated, done.



Fix it

Something's off and you can name it. Say exactly what's wrong in a follow-up: "cut the stat I never gave you," "make it shorter." Being specific about what's wrong beats starting over.



Walk away

Wrong tool, or stakes too high. Do it yourself, or take it to a person who actually knows: a teacher, a doctor, a pro.