WORK WITH AI

Context Window

When you type 2 + 2 on the calculator on your phone, you always get 4. Same input, same answer, every time. AI doesn't work that way. Consider the following: identical questions, asked to AI by two different people.

Luke

Asks AI
What car should I buy after I graduate from college?

ChatGPT answers
Great question, Luke. I definitely recommend a **Jeep Cherokee**.

Nate

Asks AI
What car should I buy after I graduate from college?

ChatGPT answers
I'm happy you asked this question, Nate. At this point in your life, I recommend the **Ford Raptor**.

Wait a second. In the last lab, you changed the answer by changing the prompt. Luke and Nate typed the exact same prompt into the same app, and still got different answers!

Some people think this is a reason not to trust AI. But the opposite is true. It's actually demonstrating the power of AI. It gave a different, more tailored answer to each.

WHY DID THIS HAPPEN?

When you use a specific model of ChatGPT, it's the same tool everyone else starts with. But ChatGPT knows more than just the last message Luke and Nate typed. For example, earlier in his chat, Nate said he loves pickup trucks. When ChatGPT answers his question, it recognizes this and suggests the Ford Raptor.

What do you call everything the model can see when it answers your question? The **context window.** Think of it almost like the working memory the model uses to answer. And it builds that working memory from several places.

WHAT THE MODEL ACTUALLY SEES

Here's the part that surprises people. When you hit send, the model doesn't just see your last message. It sees the entire context window, and it answers from all of it. The illustration above shows everything that feeds it: your prompt, everything earlier in the chat, your personalization, and saved memory. And when you work inside a project, its files and instructions ride along too.

WHAT FOLLOWS YOU WHERE

The feeds don't all have the same reach. Two live and die with a single chat. Two follow you everywhere. And projects sit in between. Every major app has some version of these, and the ones beyond the chat are yours to set up:

THIS CHAT ONLY	EVERY CHAT	CHATS IN A PROJECT
Your current prompt Everything you sent this turn: what you typed, plus anything attached.	**Personalization** What you told it once: your name, what you're into, how you like answers.	**Projects** The project's files and instructions. Inside the project, always in the window. Outside it, invisible.
Everything earlier in this chat Everything you and the AI said before now. Close the chat, and it stays behind.	**Saved Memory** What the app picked up from your chats and saved on its own.	

Project chats get everything: the every-chat feeds plus the project's own.

Here's the test: open a brand-new chat tomorrow. Before you type a single word, your personalization and memory are already in the window. Yesterday's conversation isn't, unless the app saved a piece of it to memory.

One thing to keep straight: these are app features, not the model learning about you. The app stores what you set, then quietly adds it to the context window every time you hit send.

THE WINDOW HAS A LIMIT

One more thing worth knowing: the window isn't infinite. In a long chat, the oldest parts eventually fall out. Some apps summarize them, others just forget. That's one reason starting a fresh chat for a new task is often the right move.

Big or small, what fills the window is what shapes the answer. Same question, different windows.