

Prediction

You just watched the final vector land on the map, closest to the best-matching tokens. Now here’s how the model turns that position into an actual next word.

THE RANKED LIST

The model doesn’t just grab the single closest token. It scores **every** possible next token at once and lines them up by probability: a ranked list, with the best fit on top and a long tail of less likely options below it.

Then it takes one, almost always from the top, adds it to the text, and runs the whole thing again for the token after that. That pick has a name: **prediction**. Score every possible next token, choose one, then repeat.

ONE DIAL ON THE PICK: TEMPERATURE

There’s a control on how boldly the model commits to the top of that ranked list: **temperature**.

Temperature is the variation dial. It changes how predictable or surprising the model is when choosing words. Low temperature makes answers more repeatable. High temperature makes answers more varied, creative, and sometimes weird.

Temperature does not make the model smarter. It does not make an answer true. It changes how much the model sticks to its most likely next choice.



Low temperature
Repeatable, predictable, less varied.



Medium temperature
Balanced, natural. Usually best for normal writing and brainstorming.



High temperature
Surprising, unusual, sometimes weird or incoherent.

How temperature reshapes the picks

Here’s that ranked list for “See you ___”. Watch how the same candidates redistribute as temperature rises. This is a simplified demo: real models use other sampling controls too, and exact numbers vary by model.

	 Low	 Medium	 High
tomorrow	100%	65%	40%
later	0%	20%	28%
soon	0%	10%	17%
there	0%	3%	7%
again	0%	1%	4%
other words	0%	1%	4%

You probably won’t see this slider in regular ChatGPT, Claude, or Gemini. The app usually picks the temperature for you. But the idea still matters.

WHEN YOU WANT REPEATABILITY VS VARIETY

Most students won’t touch a temperature slider directly. But you can ask for the effect through prompting.



When you want repeatability

Ask for “the standard answer,” “the most likely interpretation,” or “your best single answer.” Good for formatting, extracting information, following a template, or getting a stable result you can rely on.



When you want variety

Ask for “ten different options,” “surprise me,” or “give me unusual angles.” Good for naming, story ideas, jokes, alternate phrasings, or creative exploration. Regenerate also helps.

WHAT TEMPERATURE IS NOT

Low temperature doesn’t mean correct. It means repeatable. A low-temp model will give you the same answer every time, even when that answer is wrong. Temperature controls variation, not accuracy.



Temperature changes the variation, never the truth. It tunes how boldly the model commits to its most likely next word: down for repeatable answers, up for surprising ones. It never makes an answer more correct, only more or less varied.

These odds didn’t come from nowhere. The model learned every one of them from billions of examples, long before you typed a word. That’s **training**.