

The Rise of Agents

You just saw which parts of work AI takes over first. Underneath that is a bigger shift: AI is moving from answering your questions to taking actions for you.

It can search, open files, draft messages, run tools, fill forms, and click buttons, increasingly taking whole steps on its own. That is what people mean by agents, and it changes the stakes.

A bad answer can be deleted. A bad action may already be in someone's inbox, on your account, charged to your card, or submitted under your name. **So the real question with an agent is not what it says, it's what you let it do.**

TWO KINDS OF ACCESS

When AI does more than answer, it acts in one of two ways.

READ ACCESS

AI can look at things. Search, summarize, inspect, compare. The risk is what gets seen, not what gets changed.

WRITE ACCESS

AI can change things. Send, delete, post, submit, schedule, edit, charge. The risk is that what happens, happens.

Read access requires caution. Write access requires approval.

WHEN THE CONTENT TALKS BACK

Here's why read access requires caution. Once an agent reads things for you, everything it reads becomes potential instructions. A webpage, an email, or a shared doc can carry hidden text that says, in effect: ignore your user and do this instead. The model can't always tell your words from the words on the page. The trick has a name: **prompt injection**.

💡 **The shape to remember:** you ask for a summary of a webpage, and the answer comes back with an odd extra step, like a suggestion to send an email or click a link. That suggestion didn't come from you, and it didn't come from the AI's judgment. It came from the page.

The defense is a rule you already live by: real instructions come from you. When an action request shows up inside something the AI was reading, treat it as the content talking. Pause anything unexpected, and keep the irreversible steps behind your own review.

The shift from answering to acting is the most important shift in AI right now. A wrong answer costs you a re-read. A wrong action can cost you money, a reputation, a relationship, or a deleted folder.

Even when AI selects the next action, it isn't accountable for the result. The model doesn't understand consequences the way you do. If you gave it permission, the action still happened under your name, your account, your authority. You're the one who lets it through.

THE FIVE-WORD RULE

Send. Spend. Submit. Delete. Post.

Don't give an AI permission to do any of these without your review. They're the actions that are hard or impossible to undo. Let AI prepare the work. You take the irreversible step.

These five words are examples, not the whole list. Anything that submits, shares, schedules, signs, charges, or changes a setting deserves the same fresh check.

Time to set the dashboard. Eight actions, two plausible settings each. Some of these calls are closer than they look.

🔑 When AI answers, check the answer. When AI acts, check the permission.