

What You Learned

WHERE DOES THIS LEAVE YOU?

You see how it works.
You catch what’s wrong.
You make the call.
— *now you’re smarter than the tool.*

Everything you just learned ties back to three ideas. Lock these in and you’ve got the mental model that runs every AI conversation.



Patterns

You get why AI works at all. It read billions of pages and now spots patterns the way you spot your friend’s handwriting. Better prompts = better patterns to match.



Probability & Prediction

Every answer is built one piece at a time from probabilities. That’s why the same prompt twice can give two different answers, why fluent output can still be wrong, and why confident sounding isn’t the same as right.



Human Judgment

You set the goal, give the context, check the claims, protect what’s private, and make the final call. AI can help with the work. It can’t own any of that for you.

QUESTIONS YOU CAN ANSWER NOW

Here’s the whole course in question form.

Most people who use AI every day can’t answer these. You can.

1

Why does AI sometimes make up facts that sound completely real?

2

Why does AI produce text one token at a time?

3

Why can the same question get a different answer every time?

4

Why does context change what AI gives back?

5

Why does the way you ask a question change the quality of the answer?

6

Why does AI agree with you even when you’re wrong?

7

How do you check an AI answer before trusting it?

8

What decisions still belong to you?



You know more about AI than *most people who use it.*

You can explain what tokens are, how embeddings capture meaning, and why attention matters. Most people who use ChatGPT daily can’t. That understanding is what separates someone who uses AI from someone who uses AI *well.*