

When AI Judges You

This whole course has been about you using AI well. But there's a side of AI you don't open and don't control. AI is also being used on you. Right now, today, this week.

Maybe a system already decided which posts you saw on TikTok this morning. Maybe a flagging tool ran on your last essay before your teacher did. Maybe a tool helped sort a stack of college applications before anyone read yours in detail. These systems don't ask permission. They don't announce themselves. They make decisions about you based on data you may not know they had.

Sometimes you're the user of AI. Sometimes you're the subject of AI. The skill for the rest of your life isn't just operating AI well. It's knowing what to do when AI is operating on you.



THREE MODES TO RECOGNIZE

MODE 1 · WHAT YOU SEE

Algorithms rank what lands in front of you.

TikTok's For You feed, YouTube recommendations, Instagram Explore, search results, the order of your inbox. The system shapes your day before you even tap.

MODE 2 · WHAT YOU CAN DO

Filters and flags decide whether you can post, comment, apply, or stay logged in.

Content moderation, school plagiarism flags, account suspensions, automated rejections. Lines you can hit without seeing them coming.

MODE 3 · WHAT GETS ATTACHED TO YOUR NAME

Scores and classifications follow you.

Hiring screen ratings, college admissions tools, credit scores, plagiarism scores, fraud risk flags. A number or label that decides whether a human even looks at your case.

The modes tell you what the system is doing. The other half of reading a deployment is who's in the picture. Every system has more seats than yours, and naming the seats tells you who has power, who has risk, and who gets a say.

FOUR ROLES IN EVERY DEPLOYMENT

Same system, different stakes.

The same AI looks different depending on which side of it you're on. Name the side, and you can read the deployment.

ROLE 1 · USER

The person using the tool

Usually got what the deployment was designed to deliver. Faster work, easier access, a quicker decision. Most of the benefit is theirs. A lot of public discussion of AI focuses on them.

ROLE 2 · SUBJECT

The person the output is about

Often not the user. When AI scores a college application, the school is the user and the applicant is the subject. The subject usually has the most at stake and the least control.

ROLE 3 · OPERATOR

The company or institution running it

Could be a tech company, a school district, an employer, a hospital, or a platform. Either way, it made the deployment choice and is responsible for how the tool affects people. Their incentives are not neutral. Knowing what they were optimizing for tells you most of the story.

ROLE 4 · BYSTANDER

Everyone else who's affected

Wasn't part of the design, wasn't asked, often doesn't know the system exists. Easiest group to forget. Often where the deepest costs land.

Whichever mode it is, and whichever seat you're in, you have one tool. A short list of questions to ask any system that's judging you. Memorize them.

THE FIVE QUESTIONS

Run these on any AI that's *running on you*.

1. Who is responsible?
What person, school, company, platform, or institution chose to use the system?
2. What data was used?
What information, behavior, score, document, or profile did the system rely on?
3. Can you push back?
Is there an appeal, a moderator, a teacher, an admissions office, an employer, or some other support path?
4. What happens if it's wrong?
Is the cost small, or could it affect your grade, account, application, reputation, money, or safety?
5. Who is this helping, and who is it hurting?
Who gets the benefit when this works, and who carries the cost when it doesn't?

Notice the first question is who, not what. AI systems score, rank, flag, recommend, and sometimes act on their own. But a person or an institution chose to deploy that system, set the rules, and decide what happens when it flags someone. That's where responsibility lives, and that's who you push back on.

ASK FOR THE PAPER TRAIL

If an AI system affects a decision about you, get the decision out of the black box and into a process a human can explain. Ask, in writing if needed:

- What policy applies here?
- What data, document, score, or evidence was used?
- Was there human review?
- How do I correct a mistake?
- What is the appeal path?

A school or company that won't put any of this in writing is telling you something.

🔑 **Every system optimizes for something.** AI doesn't have intentions, but a deployed system is built to optimize for something: engagement, speed, cost, applicant sorting, fewer manual reviews. TikTok's feed is tuned for watch time. A hiring screen might be tuned to match past hires. The five questions above are all downstream of one upstream question: what is this system actually chasing, and does that line up with what's good for you?

WHEN MOSTLY RIGHT ISN'T RIGHT ENOUGH

Stakes scale with reach. A 99.9%-accurate system still gets a thousand wrong for every million people it scans. Doesn't matter if it's facial recognition, fraud detection, content moderation, or hiring screens. The averages look fine. The thousand people who got the wrong answer are real, and they don't get to opt out of the consequences because the system was "right almost all the time." Also ask: accurate for whom, and wrong in which direction? A false positive and a false negative don't hurt the same people.

WHEN FAIR-ON-PAPER ISN'T FAIR

A system can claim it doesn't use race, gender, disability, or income to make decisions. That can be true and still produce unfair outcomes.

Other data the system DOES use often correlates with what it claims it doesn't. A zip code can be a proxy for race and income. A school name can be a proxy for class. A pattern in someone's name can be a proxy for ethnicity. Even small details (gaps in a resume, the speed someone fills out a form) can correlate with things that shouldn't be the basis of any decision. These are called **proxy variables**: data points that stand in for something else, even when the system never names that thing directly.

When you ask "what data was used," the follow-up question is "what does that data correlate with?" A system that uses neighborhood, school, name, and timing patterns can be making decisions based on race and class without ever naming them.

🔑 **When AI judges you, don't argue with the machine.** Find the human, ask what data was used, and learn how to appeal.